

# Introduction to Information Security

## 12. Dynamic Analysis and Fuzzing

Kihong Heo

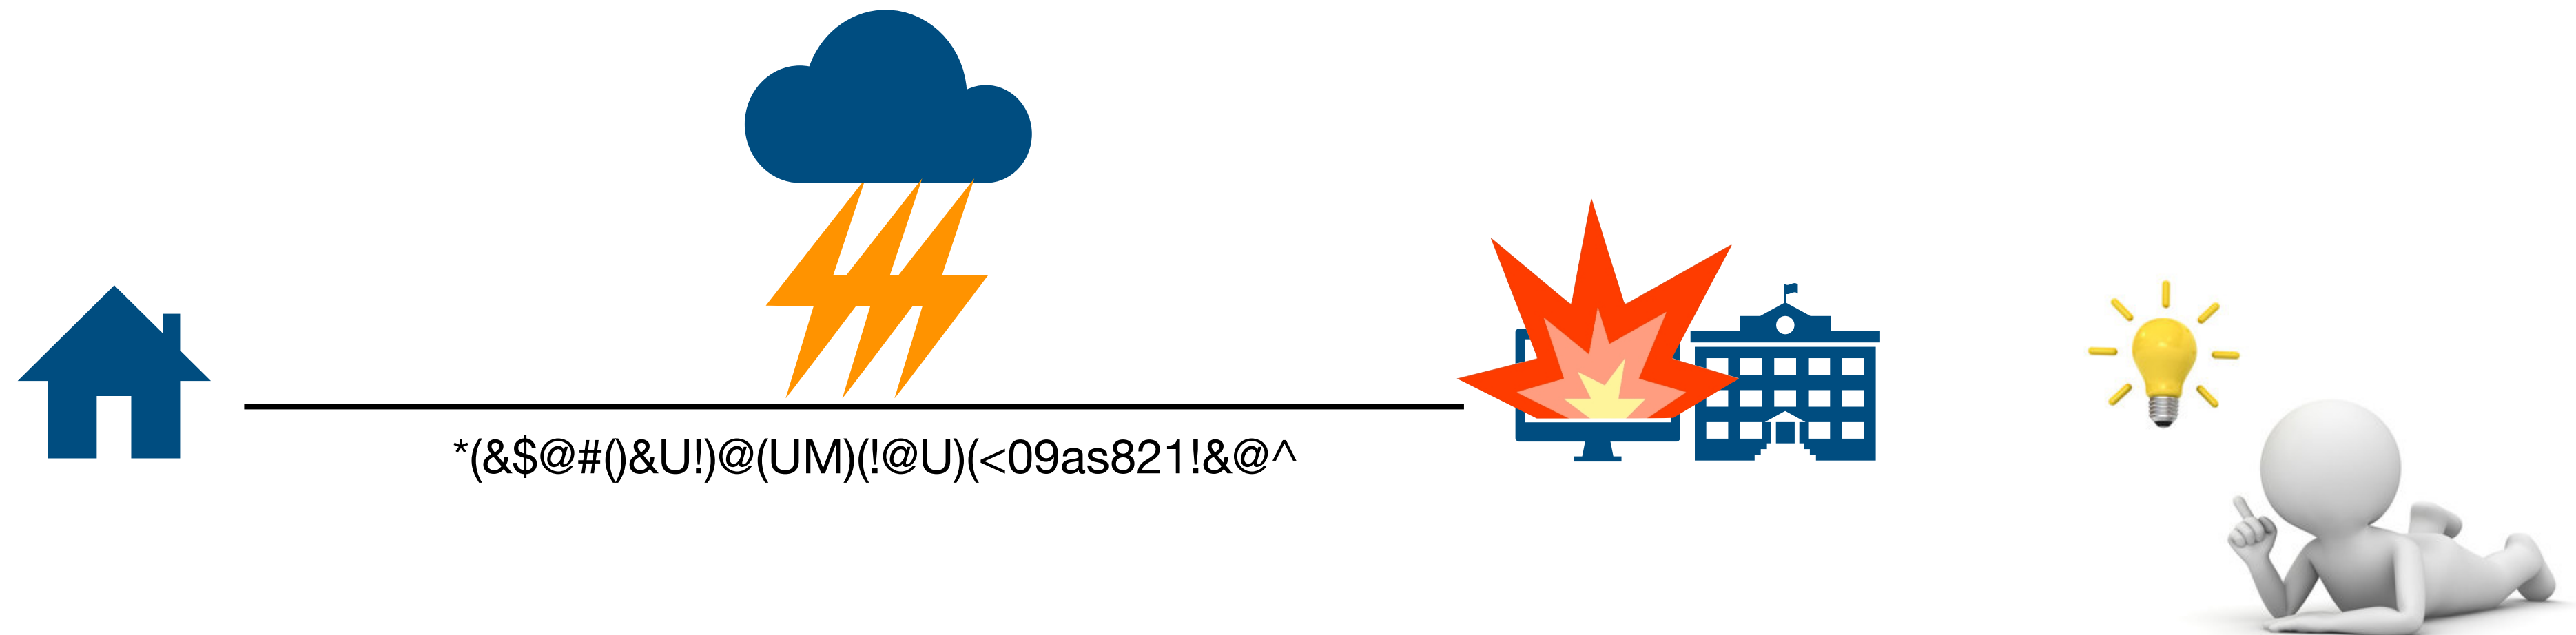


# Testing

- Check a set of **finite executions**
  - e.g., random testing (a.k.a fuzzing)
- In general, **unsound yet complete**
  - Unsound: cannot prove the absence of errors
  - Complete: produce counterexamples (i.e., erroneous inputs)
- Example: Google's oss-fuzz (<https://github.com/google/oss-fuzz>)


# Fuzzing

- Initially, developed by Barton Miller in 1988
- Thunderstorm → noise on a network line → random characters → crash
- “Can we mimic the thunder-generated noise to check robustness?”

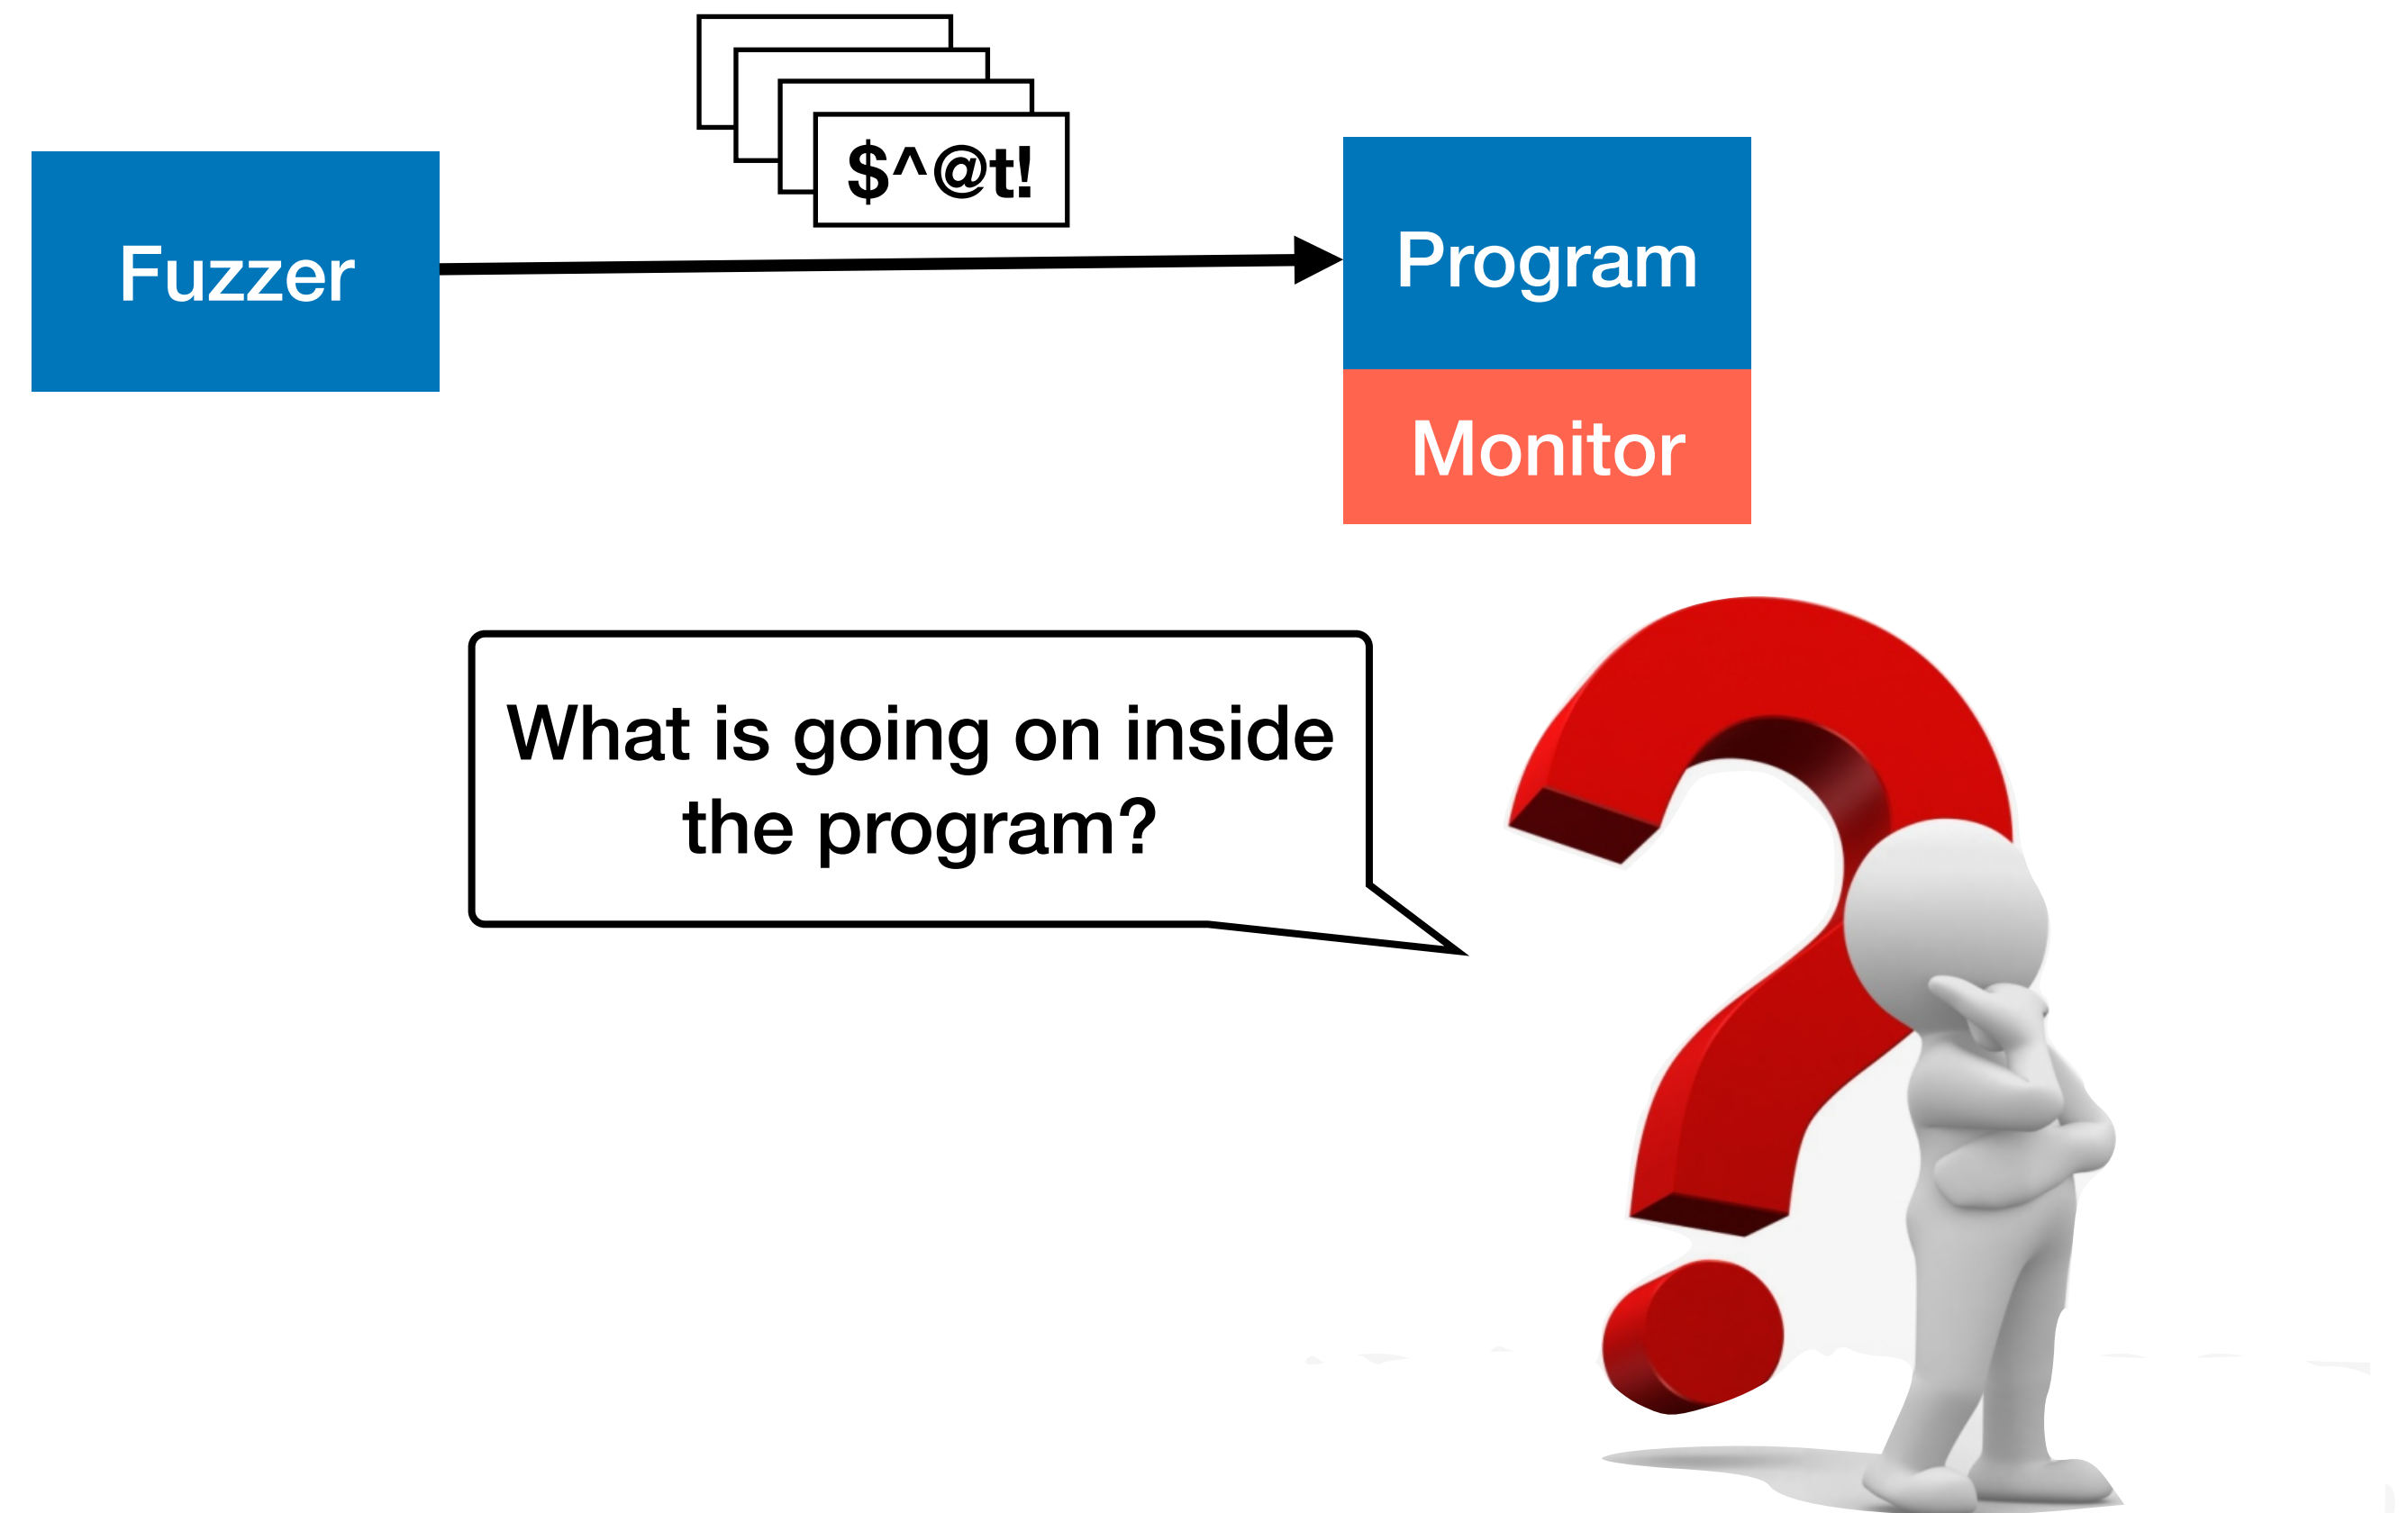


# Success Stories

- Miller et al. found many crashes in UNIX utilities
- AFL (American Fuzzy Lop) has found a lot of security vulnerabilities
  - See <https://lcamtuf.coredump.cx/afl/>
- Google's OSS-Fuzz: continuous fuzzing platform for open source SW



# Fuzzing Overview




# Runtime Monitoring

- Observe program executions
  - Explicitly trap the execution if an error happens
  - Give more detailed information (e.g., bug type, coverage, location)
- Important: achieving low overhead at runtime
- How?
  - If source code available: instrumentation via compilation(e.g., LLVM's sanitizers)
  - If no source code available: binary rewriting (e.g., Pintool) or emulation (e.g., QEMU)

# A Simple Fuzzing

- (Generate random inputs; run a given program w/ the inputs; see if it crashes)+
- Without considering the target program's detailed behavior
  - So called, **blackbox** fuzzing
- Initial success: a command-line fuzzer for UNIX utilities\*
- The infinite monkey theorem actually works!

*“A monkey randomly hitting keys for an infinite amount of time will produce any given text, such as Shakespeare’s Hamlet”*



\*Barton Miller, et al., An Empirical Study of the Reliability of UNIX Utilities, CACM, 1990



# Blackbox vs Whitebox

- Blackbox: generate inputs **regardless** of program's logic and structure

- Easy to implement and low cost

- Hard to explore deeper parts

```
x = input();  
if (x == 482716115)  
    bug();
```

- Whitebox: generate inputs by **observing** program's logic and structure

- Can explore deeper parts

- Require constraint solving (high overhead)

```
x = input();  
y = input();  
if (x < y)  
    if (x2 > 3)  
        bug();
```

```
a = input();  
b = input();  
c = input();  
n = input();  
  
if (n > 2)  
    if (an + bn == cn)  
        bug();
```

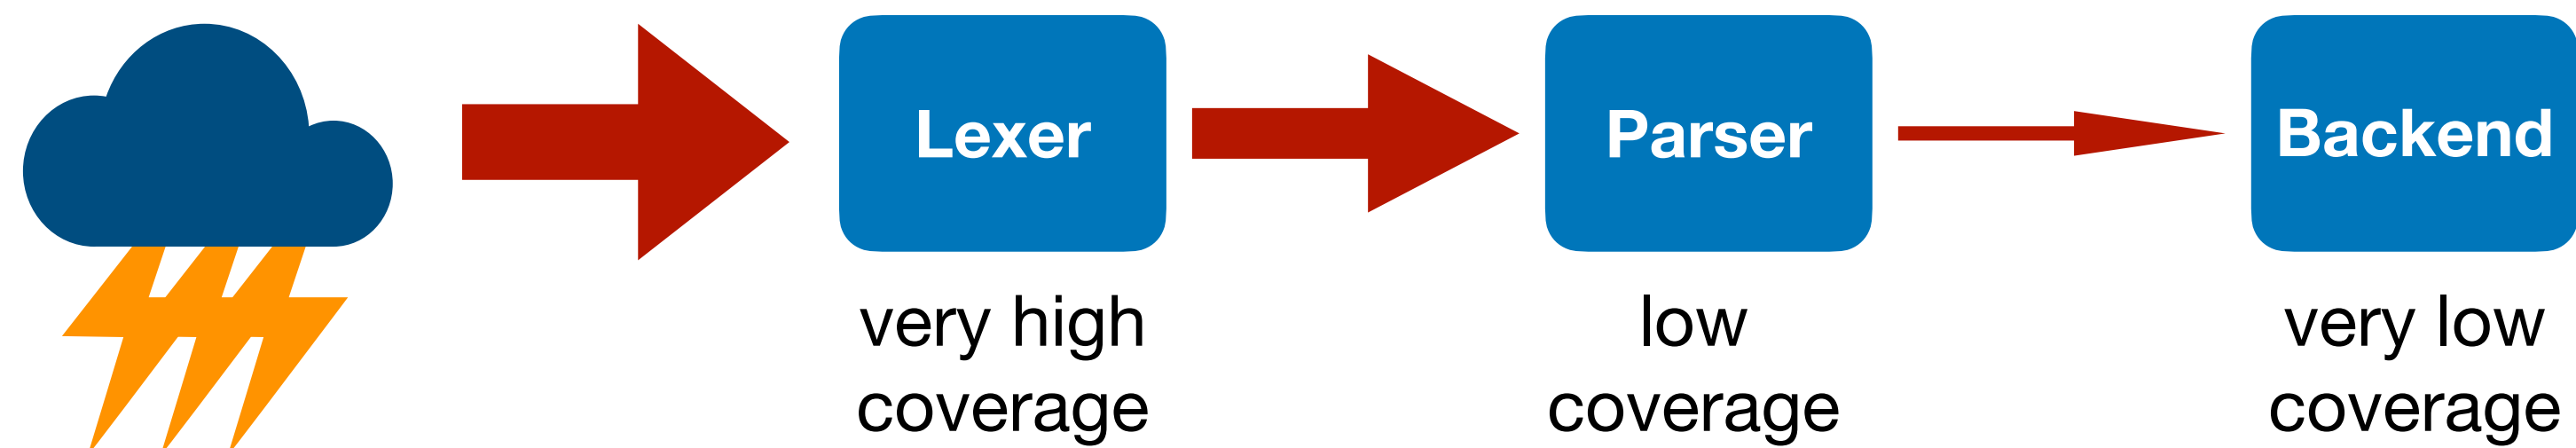


# Examples


- Monkey: a random testing tool for Android apps
  - Randomly generate streams of events such as clicks, touches, etc
  - A part of Android Studio
- Cuzz: a random testing tool for finding concurrency bugs
  - Randomly generate thread schedules
  - A part of Microsoft Application Verifier

# Problems

- Very low test coverage!
- Random inputs are often filtered out in earlier stages of programs
  - E.g., “Invalid syntax”
- What are the chances of getting valid URL from random strings?
  - URL format: `scheme://netloc/path?query#fragment`



# Fuzzing Overview



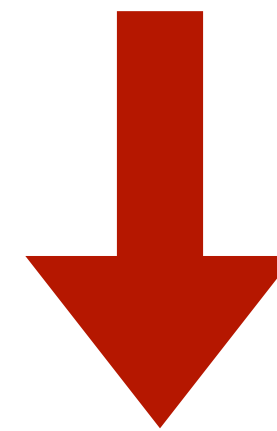
# Mutation-based Fuzzing

- Idea: generate new inputs by **mutating** existing valid inputs (seeds)
  - valid inputs usually reach deeper parts
- Mutation operators: random flips or heuristics
- Success story: AFL (2013)
  - See this blog post by the author of AFL\*

\*<https://lcamtuf.blogspot.com/2014/08/binary-fuzzing-strategies-what-works.html>

# Example: Mutating URLs

Seed input: `http://www.google.com/search?q=fuzzing`



randomly insert/delete/flip characters

0 mutations: `http://www.google.com/search?q=fuzzing`  
5 mutations: `http:/L/www.googlej.com/seaRchq=fuz:ing`  
10 mutations: `http:/L/www.ggoWglej.com/seaRchqfu:in`  
15 mutations: `http:/L/wwggoWglej.com/seaR3hqf,u:in`  
20 mutations: `htt://wwggoVgle"j.som/seaR3hqf,u:in`  
25 mutations: `htt://fwggoVgle"j.som/eaRd3hqf,u^:in`  
30 mutations: `htv://>fwggoVgle"j.qom/ea0Rd3hqf,u^:i`  
35 mutations: `htv://>fwggozVle"Bj.qom/eapRd[3hqf,u^:i`  
40 mutations: `htv://>fwgeo6zTle"Bj.\'qom/eapRd[3hqf,tu^:i`  
45 mutations: `htv://>fwgeo]6zTle"BjM.\'qom/eaR[3hqf,tu^:i`

\*<https://www.fuzzingbook.org/html/MutationFuzzer.html>

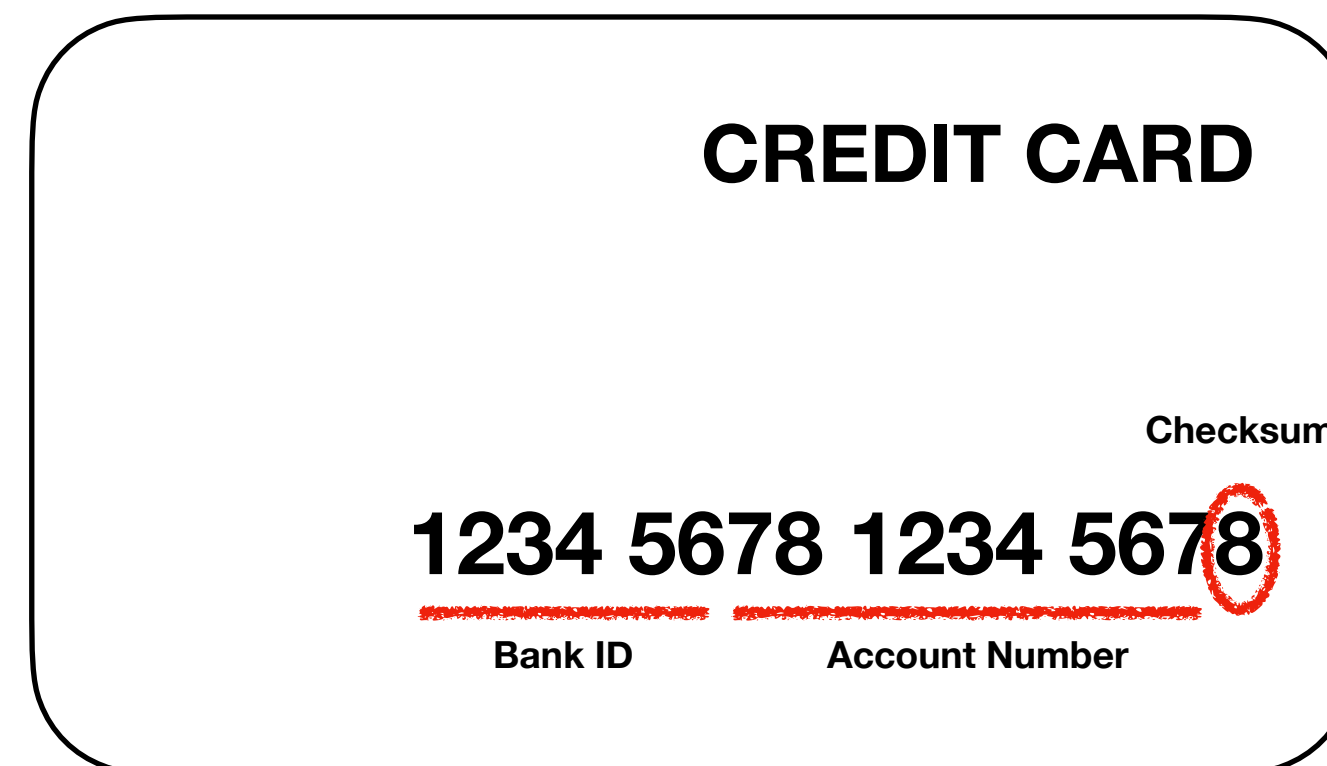
# Example Mutation Operators

- **Random bit-flipping:** randomly flip bits with a certain probability
- **Arithmetic mutation:** perform simple arithmetic on a value (e.g.,  $x + r$ )
- **Block-based mutation:** insert/delete/replace/permute/resize a subpart of an input
- **Dictionary-based mutation:** use a set of pre-defined values for mutation such as  $\{0, -1, 1\}$  or  $\{"\%s", "\%x"\}$

\*Manes et al., The Art, Science, and Engineering of Fuzzing: A Survey, TSE, 2019

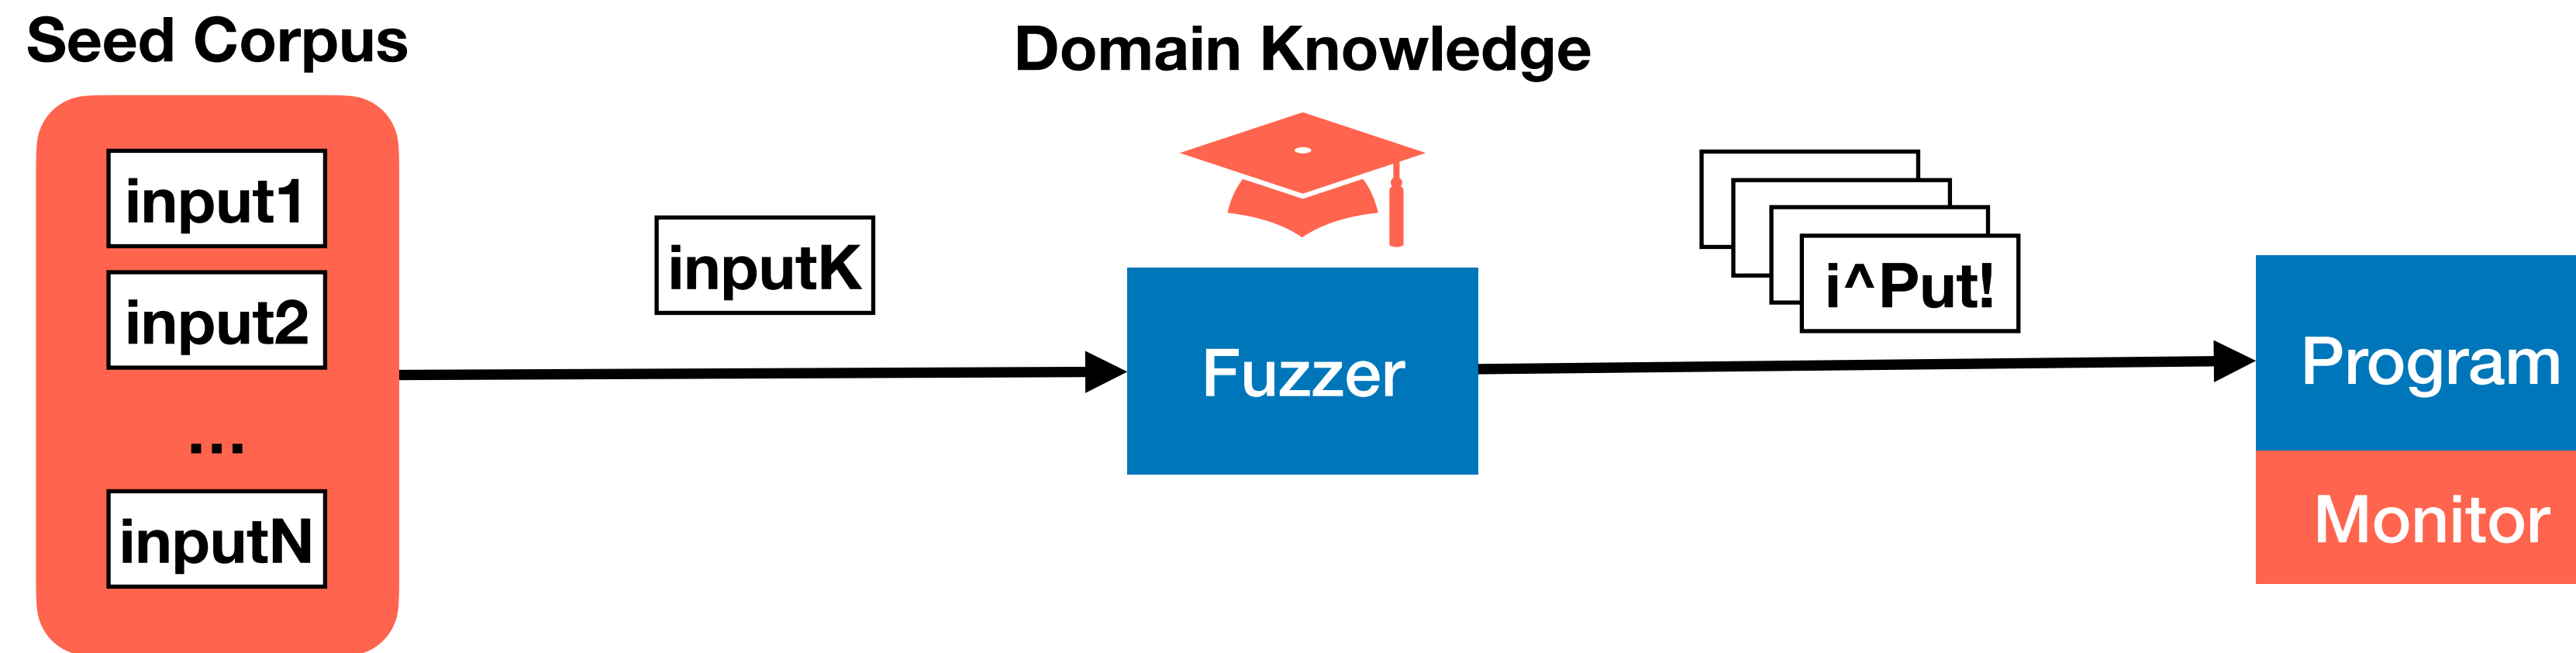
# Problems

- Limited by seed inputs
- Random mutations often violate complicated syntactic or semantics rules
  - E.g., XML, Javascript, magic number, checksum





# Fuzzing Overview



# Generation-based Fuzzing

- Generate inputs based on a **given model**
  - manually defined or automatically inferred
- Examples:
  - Kernel APIs: system call templates (i.e., function signatures)
  - Formatted data: DOM objects, PNG, MP3, etc
  - Programs: Javascript, PHP, C, etc
  - Network protocols: TLS, NFC, etc



# Problem 1: Still in the Dark

- How much fuzzing is enough?
- How to evaluate fuzzer's performance?
  - Are we going in the right direction?
  - Or stuck at some point?
- Which mutant is better than others?
  - Does it explore new execution paths?
  - Or revisit redundant paths?
- Solution: gradual, continuous, fine-grained feedback




# Problem 2: Monotonous Mutants

- Same set of corpus + same set of mutation operators = ??
- How to achieve divergence?
- Idea: divergence and selection





# Fuzzing Overview



# Coverage-guided Fuzzing

- Fuzzing as a **genetic algorithm**
  - Chromosome population: seed corpus
  - Genetic mutation: mutation
  - Fitness function: coverage
- Key idea: keep mutants that increase **code coverage** for future mutations
  - So-called **grey-box** fuzzing
  - E.g., AFL, LLVM's libFuzzer



```
arr = input();
if (arr[0] > 0) {
    if (arr[1] > 10) {
        if (arr[2] > 100) {
            bug();
        }
    }
}
```


# Code Coverage

- A metric that determines how much code has been executed
- Obtained from the runtime monitor
  - E.g., LLVM's SanitizerCoverage, gcov, etc
- Many criteria: line/stmt coverage, branch coverage, path coverage, etc
- Caveat: 100% coverage does not mean exhausted exploration

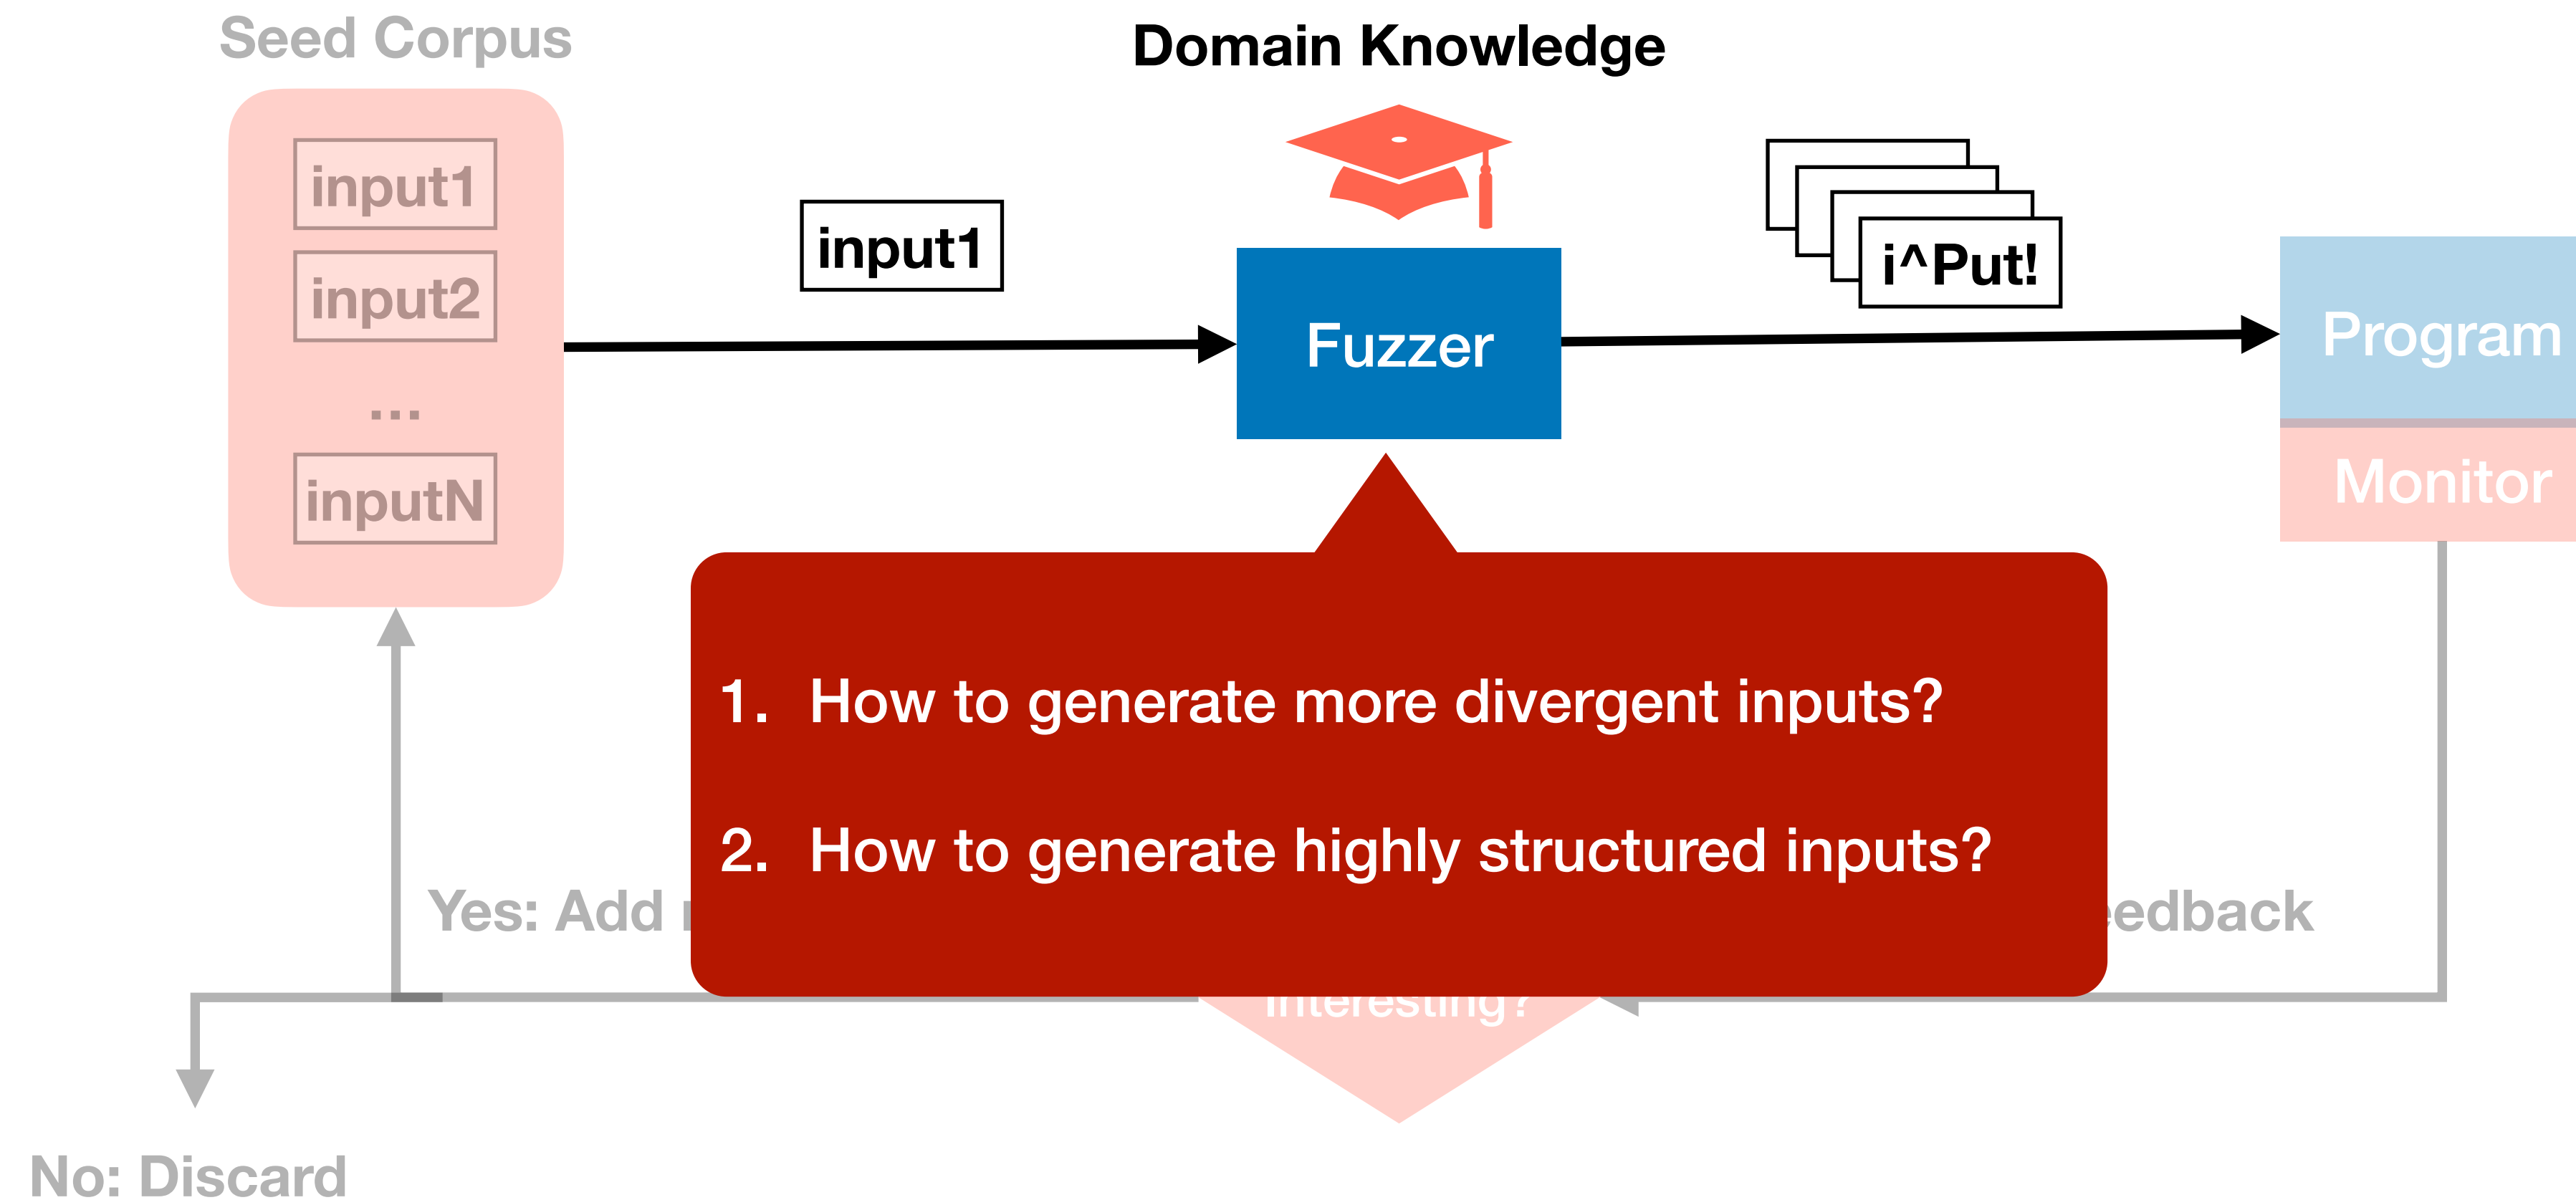
```
// x : -1, 1
x = inputs();
if (x < 0) {
    skip;
} else {
    skip;
}
assert(x != 0);
```



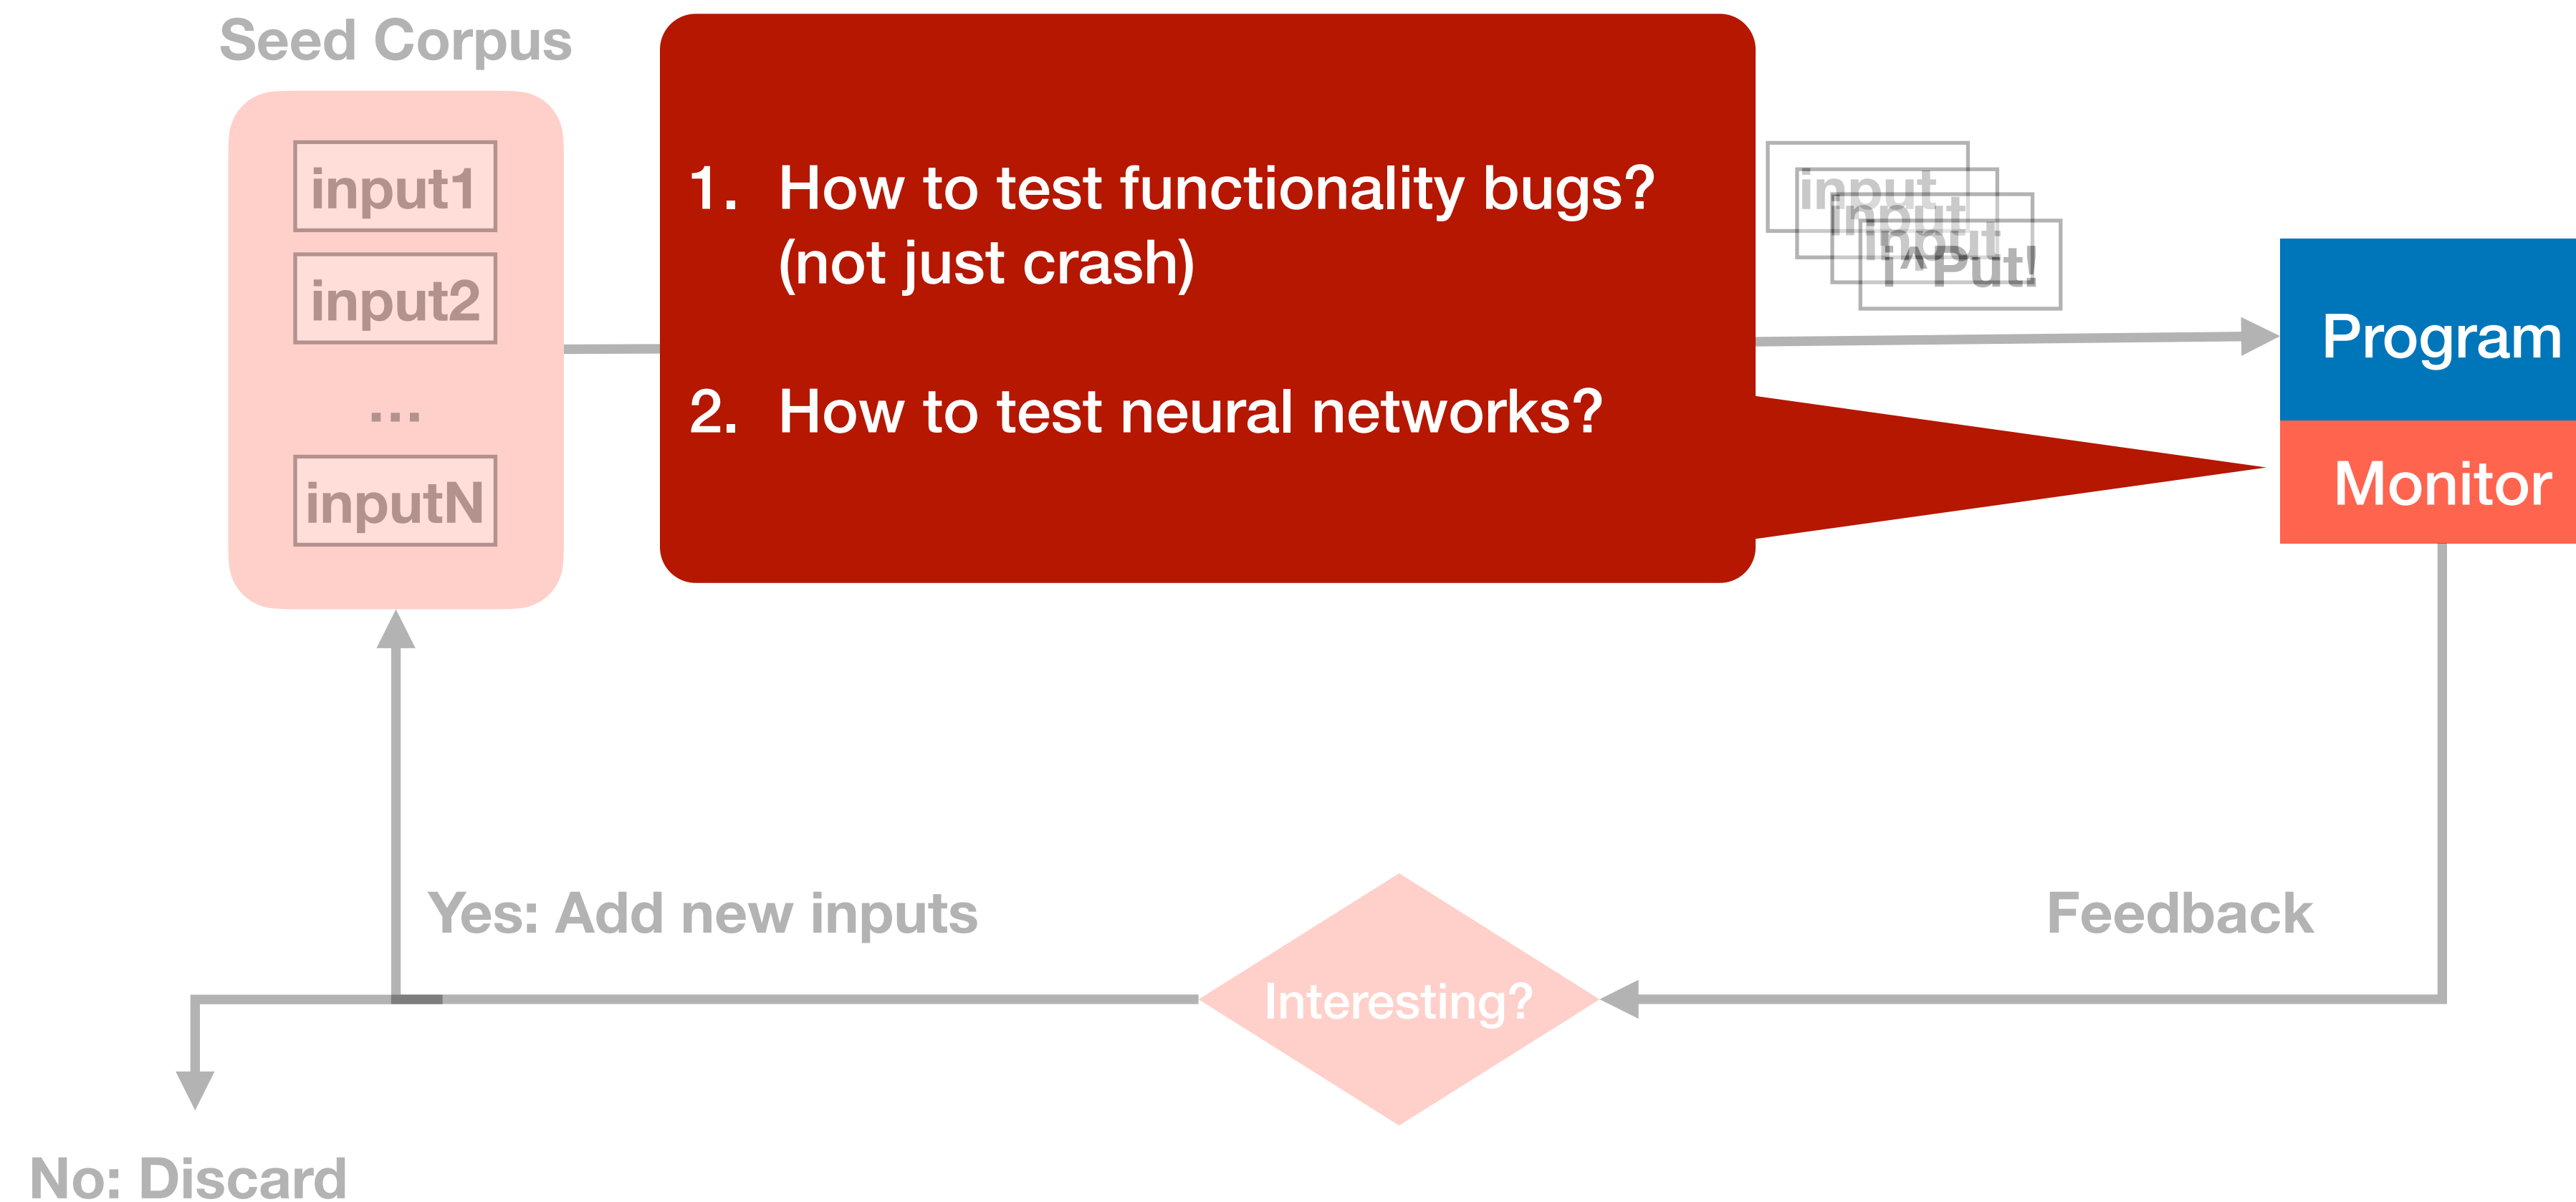
# More Research Questions



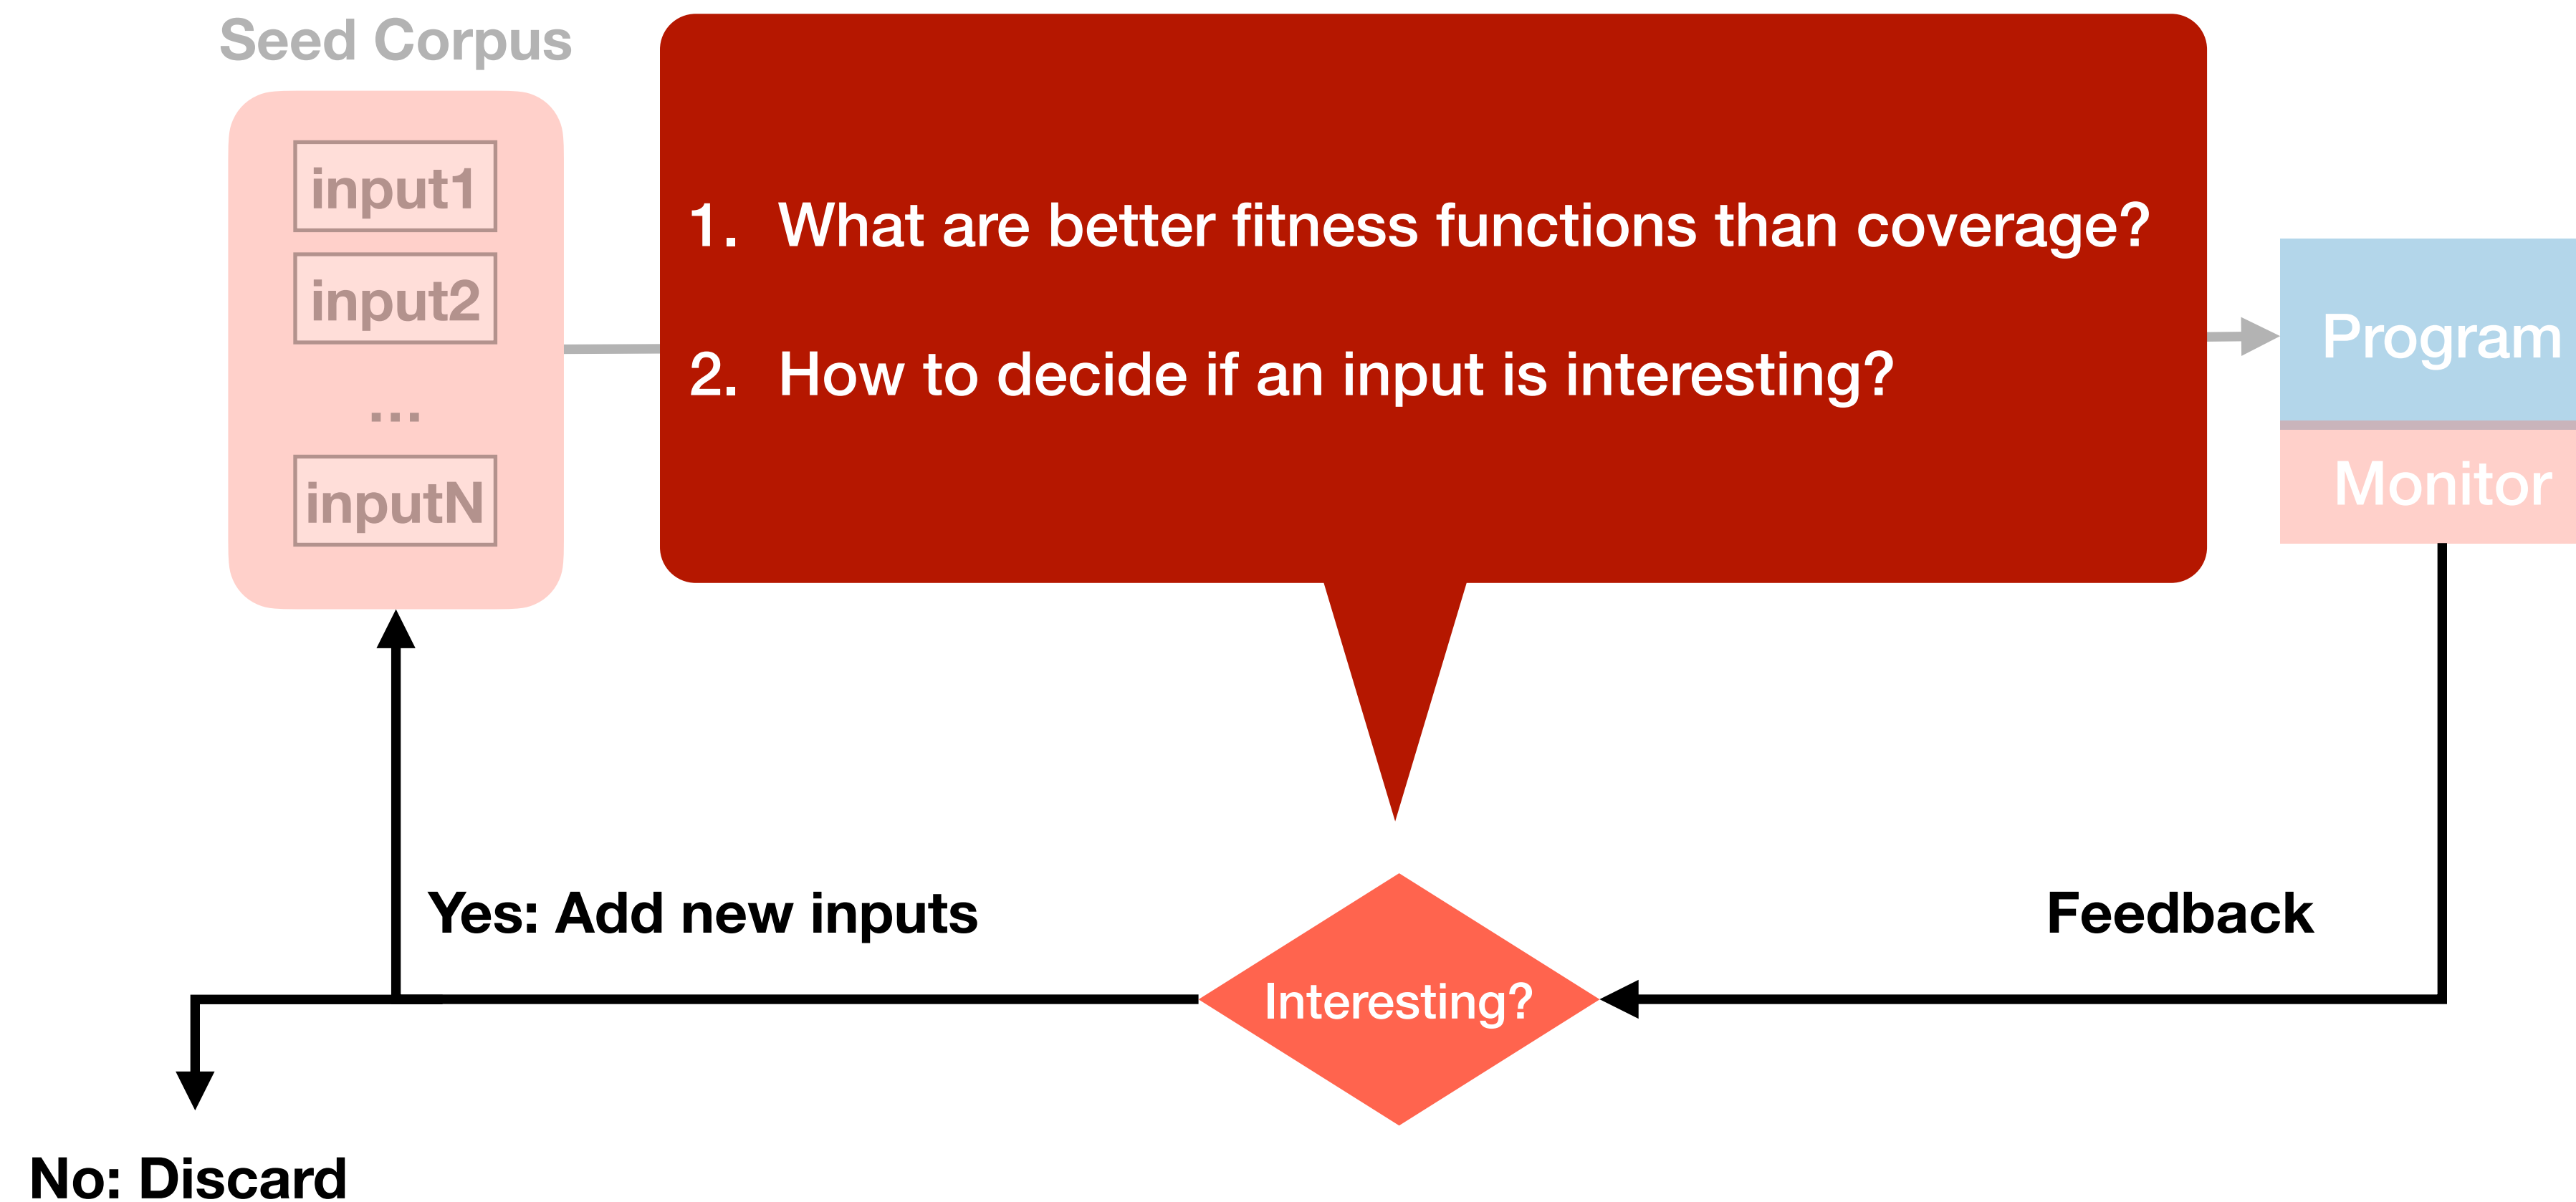
# More Research Questions



# More Research Questions



# More Research Questions



# Summary

- Fuzzing: efficient and effective testing technique
  - Input generation: mutation-based, generation-based
  - Feedback: blackbox, greybox, whitebox
- Complete but unsound with respect to program correctness
- Challenges
  - Efficiency (e.g., higher coverage)
  - Expressiveness (e.g., functionality errors)
  - Etc